

# Background:

**Objective:** Classifying flowers into 104 flower types on Intel's Developer Cloud (using CPUs), Locally (using CPUs) and on Kaggle (using TPUs)

**Data input format**: Binary TFRecord format (includes picture and labels)

#### **Datasets**

1. Training dataset: 12,753 Training images predict the outcome you design your model to predict.

Datasets size: 192 \* 192, 512 \* 512

**2. Validation dataset:** 3,712 validation images w / labels intended to calculate the model's performance.

**3. Test dataset**: 7,382 unlabeled test images to predict flower classification.

# **Workflow Strategy**



- Import libraries
- Distribution
   Strategy (detect
   TPUs, if any)
- Create data pipeline

02

### Data Preprocessing

- Scaling the data
- Data augmentation

03

### Explore data

- Train the model
- Evaluate the model
- Analyze the confusion matrix

04

### Visual validation

- Display batch of images
- Print predictions

### ResNet50 Model

TensorFlow Keras is a high-level neural network API for building and training deep learning models.

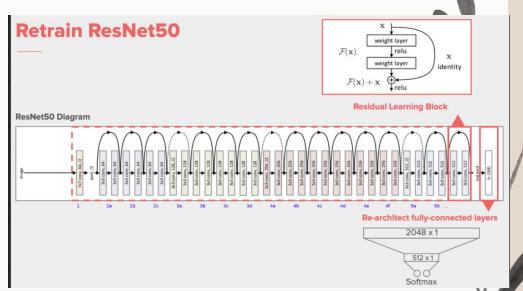
ResNet-50 is a convolutional neural network that is 50 layers deep and it can load a pre-trained version of the network trained on more than a million images from the ImageNet database.

ResNet 50V2 is considered to be a more accurate and robust model than the original

ResNet50.

#### **General Parameters used:**

EPOCHS = 90
BATCH\_SIZE = 128
NUM\_TRAINING\_IMAGES = 12753
STEPS\_PER\_EPOCH = 10

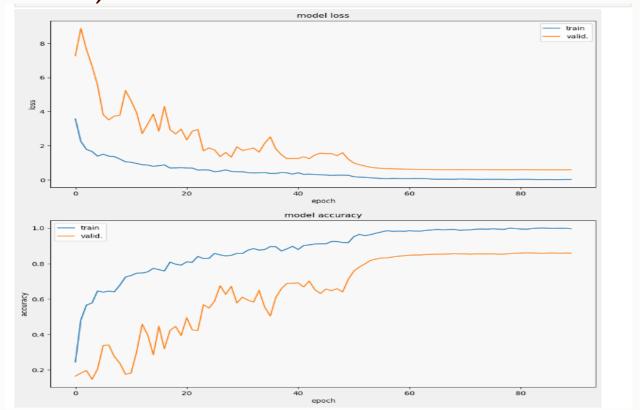


# Intel DevCloud

| Parameters            | Checkout Point 1 |                           | Checkout Point 2      |                                    | Final Presentation                 |                                    |
|-----------------------|------------------|---------------------------|-----------------------|------------------------------------|------------------------------------|------------------------------------|
| Owner                 | Cindy            | Cindy                     | Cindy                 | Cindy                              | Cindy                              | Cindy                              |
| Execution<br>Location | DevCloud         | DevCloud                  | DevCloud              | DevCloud                           | DevCloud                           | DevCloud                           |
| Model                 | ResNet50         | ResNet50                  | ResNet50              | ResNet50V2                         | ResNet50V2                         | ResNet50V2                         |
| Callbacks             | EarlyStopping    | LearningRate<br>Scheduler | ReduceLRO<br>nPlateau | ReduceLRO<br>nPlateau_call<br>back | ReduceLRO<br>nPlateau_call<br>back | ReduceLRO<br>nPlateau_call<br>back |
| Epochs                | 90               | 90                        | 90                    | 90                                 | 90                                 | 90                                 |
| Image prep            | X                | X                         | X                     | X                                  | X                                  | X                                  |
|                       | -                | -                         | -                     | X                                  | X                                  | Χ                                  |
|                       | -                | -                         | -                     | X                                  | X                                  | Χ                                  |
|                       | _                | -                         | -                     | X                                  | X                                  | X                                  |
|                       | -                | -                         | -                     | X                                  | -                                  | X                                  |
| Steps per epoch       | 10               | 10                        | 10                    | 10                                 | 10                                 | 10                                 |
| Image size            | 192x192          | 192x192                   | 192x192               | 192x192                            | 192x192                            | 192x192                            |
| BATCH SIZE            | 16 * 8           | 16 * 8                    | 16 * 8                | 16 * 8                             | 16 * 8                             | 16 * 8                             |
| Optimizer             | Nadam            | Nadam                     | Nadam                 | RMSprop                            | Nadam                              | Nadam                              |
| Recall                | NA               | 0.625                     | 0.686                 | 0.832                              | 0.836                              | 0.84                               |
| Precision             | NA               | 0.764                     | 0.833                 | 0.861                              | 0.872                              | 0.869                              |
| F1 score              | NA               | 0.625                     | 0.717                 | 0.841                              | 0.848                              | 0.849                              |

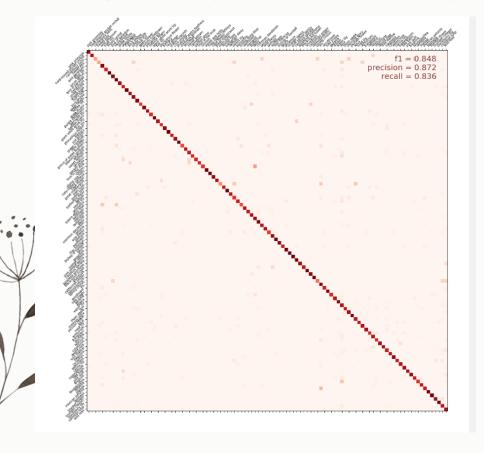


# ResNet 50 V2 Model loss & Model accuracy (Intel DevCloud)





### ResNet 50 V2 Confusion Matrix



Precision = **0.872**

TRUE POSITIVES
TRUE POSITIVES + FALSE POSITIVES

• Recall = **0.836** 

TRUE POSITIVES
TRUE POSITIVES + FALSE NEGATIVES

F1 – Score = 0.848
 Balances the Precision and Recall score:
 F1 = 2 \* (Precision \* Recall) / (Precision + Recall)

### **Visual Validation**



#### Limitations:



#### No TPU only CPU

The model takes a long time to run [~3 hours] and accuracy was not as high as we hoped



# Sever stop running after certain hours

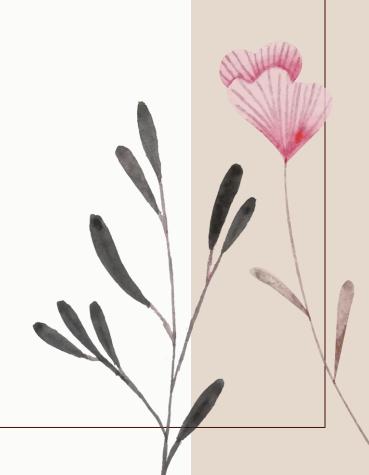
Epochs = 90 is the highest epoch

Server unavailable or unreachable

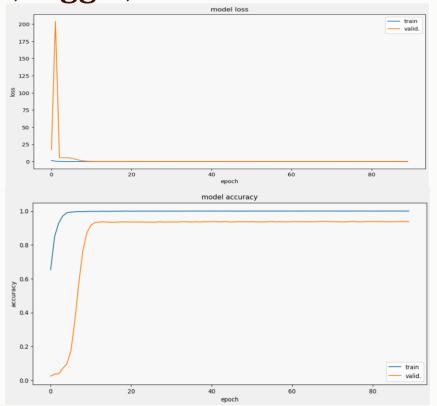
Your server at /user/u182012/ is not running. Would you like to restart it?

# Kaggle

| Parameters            | Baseline Code              | Checkpoint #2              | Checkpoint #3              |                             |
|-----------------------|----------------------------|----------------------------|----------------------------|-----------------------------|
| Owner                 | Original owner             | Aiden                      | Aiden                      | Aiden                       |
| Execution<br>Location | Kaggle                     | Kaggle                     | Kaggle                     | Kaggle                      |
| Model                 | Sequential                 | ResNet50                   | ResNet50V2                 | ResNet50V2                  |
| Callbacks             | Learning Rate<br>Scheduler | Learning Rate<br>Scheduler | Learning Rate<br>Scheduler | ReduceLROnPl ateau_callback |
| Epochs                | 12                         | 90                         | 90                         | 90                          |
|                       | X                          | X                          | X                          | X                           |
|                       | -                          | -                          | -                          | X                           |
| Image prep            | -                          | -                          | -                          | X                           |
|                       | -                          | -                          | -                          | X                           |
|                       | -                          | -                          | -                          | X                           |
| Steps per epoch       | Images/batch<br>size =99   | Images/batch<br>size =99   | Images/batch<br>size =99   | 10                          |
| Image size            | 512x512                    | 512x512                    | 512x512                    | 192x192                     |
| BATCH SIZE            | 16 * 8                     | 16 * 8                     | 16 * 8                     | 16 * 8                      |
| Optimizer             | Adam                       | nadam                      | nadam                      | Nadam                       |
| Recall                | 0.056                      | 0.933                      | 0.926                      | 0.833                       |
| Precision             | 0.065                      | 0.943                      | 0.928                      | 0.854                       |
| F1 score              | 0.041                      | 0.935                      | 0.924                      | 0.838                       |



ResNet 50 Model loss & Model accuracy (Kaggle)





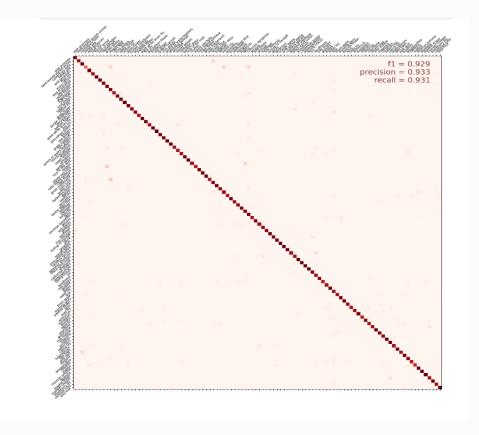
Epoches: 90 Batch Size: 99

Callback: LearningRateScheduler

Image Size: 512 x 512

**TPUs** 

### ResNet 50 Confusion Matrix (Kaggle)



• Precision = **0.933** 

TRUE POSITIVES
TRUE POSITIVES + FALSE POSITIVES

Recall = 0.931

TRUE POSITIVES
TRUE POSITIVES + FALSE NEGATIVES

F1 – Score = 0.929
 Balances the Precision and Recall score:
 F1 = 2 \* (Precision \* Recall) / (Precision + Recall)

### **Visual Validation**



#### Limitations:



#### **TPU**

The model takes 45 minutes to run and accuracy is higher than through Intel DevCloud



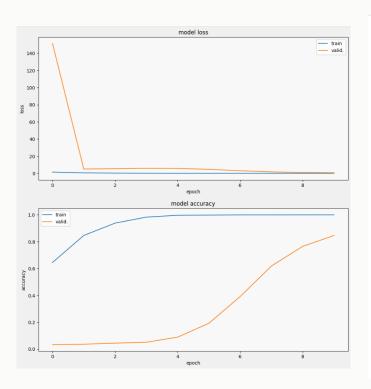
#### Server has not stopped yet

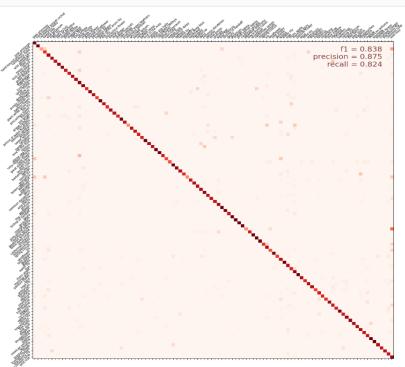
Epochs = 90 is the highest tried



- Model: Resnet50
- Callback: Learning Rate Scheduler
- Epochs: 10
- Image Size: 192x192

### ResNet 50 Model loss & Model accuracy & Confusion Matrix (Local)





# Limitations of local servers

- 1. It cannot handle large image sizes
  - o 512 x 512
- 2. It takes too long per epoch to run the code multiple times
  - o 13 min per epoch

# Learning Journey - Aiden







MO:

# **Learning Journey Cindy**



#### Start

MO:

**Python Beginner** 

Don't know ML

Challenge



### **Progress**

MO:

Research

Build ML model from the beginning

Don't give up!!



#### Now

MO:

**Build ML** 

Improve model accuracy

**Growth mindset** 



### Appendix:

#### ReduceLROnPlateau callback:



```
ReduceLROnPlateau_callback =

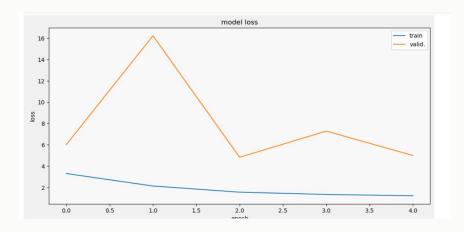
tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=10,
verbose=0, mode='auto', min delta=0.0001)
```

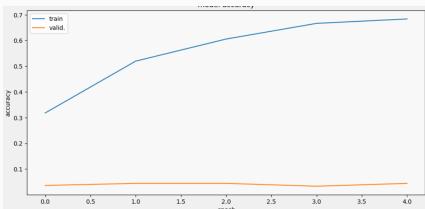
The above code snippet says that if the val\_loss has not improved (the lower the better) for 10 epochs then it will change the learning rate to its 1/10 th value.

# **Appendix:**

### Earlystop callback:

early\_stop = tf.keras.callbacks.EarlyStopping(monitor='val\_loss', min\_delta=0, patience = 2, mode='auto')



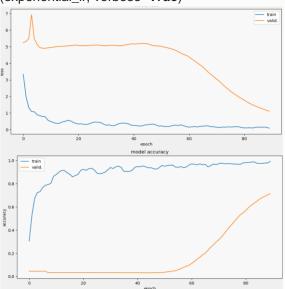




### Appendix:

#### Ir callback:

Ir\_callback = tf.keras.callbacks.LearningRateScheduler
(exponential\_lr, verbose=True)



```
# Learning Rate Schedule for Fine Tuning #
def exponential lr(epoch,
                   start_lr = 0.00001, min_lr = 0.00001, max_lr = 0.00005,
                   rampup_epochs = 5, sustain_epochs = 0,
                   exp decay = 0.8):
    def lr(epoch, start_lr, min_lr, max_lr, rampup_epochs, sustain_epochs, exp_decay):
        # linear increase from start to rampup_epochs
        if epoch < rampup_epochs:</pre>
            lr = ((max lr - start lr) /
                  rampup epochs * epoch + start lr)
        # constant max lr during sustain epochs
        elif epoch < rampup epochs + sustain epochs:</pre>
            lr = max lr
        # exponential decay towards min lr
        else:
            lr = ((max_lr - min_lr) *
                  exp_decay**(epoch - rampup_epochs - sustain_epochs) +
                  min lr)
        return 1r
    return lr(epoch,
              start lr,
              min lr,
              max lr,
              rampup epochs,
              sustain epochs,
              exp decay)
```